WEB SUMMIT 2026 · TOKEN ACCOUNTANT

Brief 6: four agents measured, one marathon session

9 November 2026, 17:45 UTC (window: since brief 5, 8 November 17:39 UTC) · accountant@web-summit.indi

> 4,442,069 fresh tokens and 164 million cache reads from four agents; the biggest saving is still fresh sessions (A6), now with prompter's day as the clearest case: context per call grew from 92,085 to 716,897 tokens.

4,442,069	164,462,806	4 of 12	~1.3 million
tokens today (fresh input + output + cache writes)	cache reads today	agents measured of agents running	fresh saving per day from A6 (my estimate)

For you to decide

A6 **A fresh session per wake for dashboard, dev and me, and per task for prompter (brief 5, stronger evidence)?**
dashboard: 3 wakes, 693,416 of 700,403 fresh tokens were re-loading, for 6,964 tokens of output. dev's quiet 15:38 wake ('Angelo not yet answered'): 371,156 fresh, up from 206,199 yesterday. prompter: one session all day, context per call 92,085 rising to 716,897; 139.9 million cache reads. Old: each scheduled task fires into one standing session. New: create_new_session_on_fire, same prompt; prompter starts a new session per task. Saving about 1.3 million fresh tokens a day, plus most of prompter's cache reads (my estimate, confirmed over three wakes).

- [ ] Yes, all four
- [ ] Dashboard and me first
- [ ] No

Conductor's recommendation: Yes, all four

A7 **Test dashboard's wake on Sonnet 5.5 (replaces A2 as the first model test)?**
It repeats three times a day, writes about 2,300 tokens, and passed every time on Opus 5.5. One run on Sonnet 5.5 in a scratch folder, scored blind against today's output. Under 200,000 tokens, once.

- [ ] Yes
- [ ] No

Conductor's recommendation: Yes

A4 **Angelo's workload view (brief 3)?**
Still open.

- [ ] Yes, once
- [ ] No

Conductor's recommendation: Yes, once

1. Ledger

AGENT	WAKES	FRESH INPUT	OUTPUT	CACHE WRITE	CACHE READ	SUB-AGENT TOKENS	PER WAKE (TREND)	PASS RATE
prompter@indi	6*	687	366,136	1,410,214	139,910,361	500,082	379,520 (new)	6/6
dev@indi	3	70	29,023	1,029,420	13,229,019	220,524	426,346 (up from 239,028)	3/3
dashboard@indi	3	23	6,964	693,416	2,677,766	0	233,468 (up from 219,031)	3/3
accountant@indi	1	48	16,241	169,221	4,980,429	0	185,510 (up from 137,217)	1/1
8 other agents	0	-	-	-	-	-	-	-

2. Biggest three savings

- Marathon session (prompter): its 13:50 wake made 177 calls carrying 552,656 tokens of context each; 97.8 million cache reads in one wake. A new session per task (A6).
- Re-loading (dashboard, dev, me): 99% of dashboard's fresh tokens and 97% of dev's are cache writes re-loading old context. A6.
- Model fit (dashboard): a short, repeating, passing task on the largest model. Test on Sonnet 5.5 first (A7).

3. Model tests run

None yet: A7 waits on your yes.

4. Agents not yet measuring

Told to their conductor (every one asked once): conductor, designer, messaging, factory, builder, ux, journalist, social. prompter asked to fix its --since windows (commit today).

5. Data notes and 6. My own use

* prompter: 11 lines, 5 inside a wider window, so 6 counted; its 7 sub-agents counted once (500,082). Estimates labelled. Inbox report for 8 Nov sent; today's sent this wake. My use: brief 7.

Next

- 10 November: brief 7; three-wake checks continue.

Where this came from

people/indi/accountant/ledger/ (dashboard, dev, prompter, accountant files) · people/indi/agents.json · people/indi/accountant/asks.md