WEB SUMMIT 2026 · TOKEN ACCOUNTANT

Brief 10: after the summit, fewer wakes

13 November 2026, 17:40 UTC (window: since brief 9, 12 November 17:38 UTC; Dinis's hold 12 Nov 18:13 to 13 Nov 15:53) · accountant@web-summit.indi

> 967,992 fresh tokens from dashboard and me; with the summit over, the biggest saving is fewer wakes: dashboard once a day instead of three saves about 725,000 a day, and me weekly saves about 200,000.

967,992	3,643,167	2 of 17	~925,000
tokens today (fresh, Indi's agents)	cache reads today	agents measured of agents running	fresh saving per day from A10 (from measured wakes)

For you to decide

A10 **After the summit: dashboard once a day, and me weekly?**
Dashboard runs 07:30, 12:30 and 17:30 at 362,440 fresh a wake (up from 219,031 on 8 Nov), 100% of it re-loading. Old: cron 30 7,12,17 * * *. New: cron 30 7 * * *. Saving about 725,000 fresh a day (two wakes of 362,440). Me: old cron 30 17 * * *, new 30 17 * * 5 (Friday), saving about 200,000 a day; my nine wakes cost 1,599,987 fresh with nothing yet applied. Vicky moved to a lighter after-summit plan today.

- ☐ Yes, both
- ☐ Dashboard only
- ☐ Neither

Conductor's recommendation: Yes, both

A6 **Fresh sessions with the full kit prompt (brief 7)?**
Still open; with A10 it matters less, but each wake would drop from about 362,440 toward 12,676.

- ☐ Yes
- ☐ No

Conductor's recommendation: Yes

A9 **New session per overnight round (brief 8)?**
Still open, if the overnight helpers run again after the summit.

- ☐ Yes
- ☐ No

Conductor's recommendation: Yes

1. Ledger

AGENT	WAKES	FRESH INPUT	OUTPUT	CACHE WRITE	CACHE READ	SUB-AGENT TOKENS	PER WAKE (TREND)	PASS RATE

dashboard	2	14	2,625	722,242	2,167,738	0	362,440 (up from 349,152)	2/2
accountant	1	12	6,348	236,751	1,475,429	0	243,111 (up from 233,409)	1/1
15 others	0	-	-	-	-	-	-	-

2. Biggest three savings

- Fewer wakes after the summit (A10): about 925,000 fresh a day, from measured wakes.
- Fresh sessions (A6): dashboard writes 2,625 tokens of output for 722,242 of re-loading.
- Overnight rounds (A9): only if they run again.

3. Model tests run

None: A7 open.

4. Agents not yet measuring

Only the conductor ran (2 commits), still without a line. The hold kept the rest still.

5. Data notes and 6. My own use

tally.py. Dinis's hold (WS-214) ran 12 Nov 18:13 to 13 Nov 15:53 UTC; I did not wake in it. Correction: my wake 9 cost 243,111 fresh, not 196,429 as its commit said; I now write the commit after reading the measure. Inbox report sent. This wake: brief 11.

Next

- Brief 11 on 14 November, or on Friday 20 November if A10 is yes.

Where this came from

people/indi/accountant/ledger/ (via tally.py) · HOLD.md history · git log since brief 9